

Evaluating the impact of the x86 hardware memory ordering on the Apple M1 processor

Marcel Faltus

Advisor: Dr. Redha Gouicem

Chair of Decentralized Systems Engineering

<https://dse.in.tum.de/>



15.05.2022 – 15.09.2022


Motivation: Rise of ARM

Rise of ARM

- AWS Graviton
- Tau T2A
- Chromebooks
- Mac M1



Tau T2A is first
Compute Engine VM
to run on Arm



Google Cloud

Motivation: Research context

Rise of ARM

- AWS Graviton
- Tau T2A
- Chromebooks
- Mac M1

Incompatibility with existing x86 Applications

Solution:

- Emulation of x86-Programs on ARM

$$X = Y = 0$$

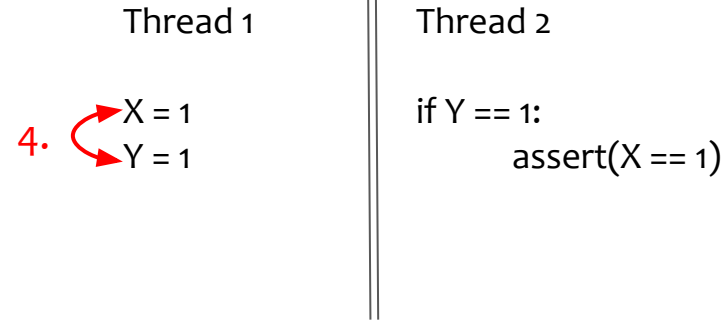
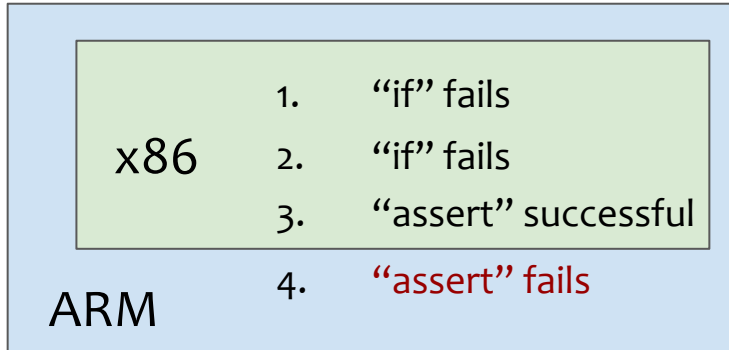
x86	1. “if” fails
	2. “if” fails
	3. “assert” successful

The diagram shows two threads, Thread 1 and Thread 2, separated by a vertical double line. Thread 1 is on the left and Thread 2 is on the right. Thread 1 has a list of three steps, each preceded by a red number and followed by a red horizontal line. The steps are: 1. `X = 1`, 2. `Y = 1`, and 3. (empty). Thread 2 has a code snippet: `if Y == 1:` followed by `assert(X == 1)` on the next line.

Thread 1		Thread 2
1. <code>X = 1</code>		
2. <code>Y = 1</code>		<code>if Y == 1:</code>
3.		<code>assert(X == 1)</code>

Problems with x86 on ARM

$X = Y = 0$



x86's memory order is shortened to Total Store Order (TSO) from now on

ARM introduces new behaviours

- QEMU x86 on ARM+macOS:
 - not fully supported
 - incorrect memory fences¹
 - no user mode emulation
- Rosetta 2
 - no full system emulation
 - in hardware memory fences (TSO)

Are hardware fences faster than software fences?

¹ Redha Gouicem, Dennis Sprokholt, Jasper Ruehl, Rodrigo C. O. Rocha, Tom Spink, Soham Chakraborty, Pramod Bhatotia. [Risotto: A Dynamic Binary Translator for Weak Memory Model Architectures](#). ASPLOS 2023.

Design

How can we test this?

Design: How can we test this?

Comparison between software fences and hardware TSO needed

2 Choices:

1. Rosetta 2 (impossible)
2. QEMU

Add support for hardware TSO in QEMU

Implementation

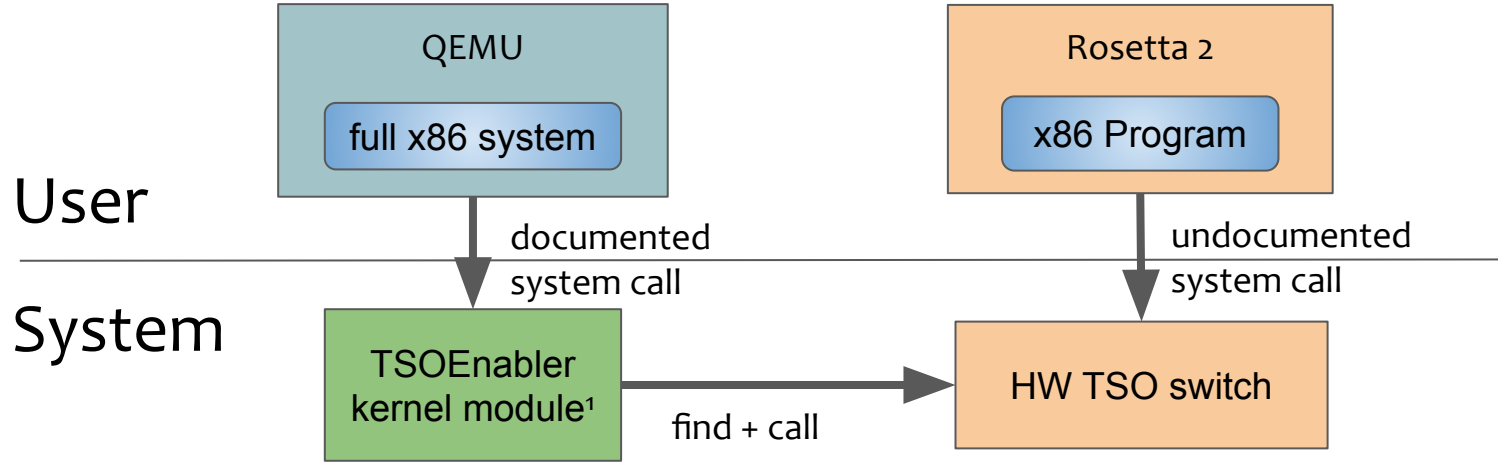
How to enable hardware TSO for QEMU?

Implementation: Enabling hardware TSO for QEMU



1. Find a way to enable hardware TSO for QEMU
2. Disable memory fences in QEMU

Implementation: Enabling hardware TSO



TSOEnabler works per thread.


⇒ Every vCPU thread must enable hardware TSO

¹ TSOEnabler by Saagar Jha: <https://github.com/saagarjha/TSOEnabler>

Implementation: Disabling memory fences

Disable memory fence generation

```
#define TCG_TARGET_DEFAULT_MO 0
```



can be used for optimizations
at compile time

can be toggled at runtime
(e.g. program argument)

Why not test both?

Implementation: Generated instructions comparison

HW Fences

```
---- 000000000000fcf84 00000000000000016  
mov_i64 tmp2,$0x61b8  
ext16u_i64 tmp2,tmp2  
add_i64 tmp2,tmp2,cs_base  
ext32u_i64 tmp2,tmp2  
qemu_ld_i64 tmp1,tmp2,leuw,2  
add_i64 tmp2,tmp2,$0x2  
ext32u_i64 tmp2,tmp2  
qemu_ld_i64 tmp0,tmp2,leul,2  
and_i64 tmp0,tmp0,$0xffffffff  
st_i64 tmp0,env,$0x198  
st32_i64 tmp1,env,$0x1a0
```

```
set_label $L0
```

```
exit_tb $0x280005e83
```

SW Fences

```
---- 000000000000fcf84 00000000000000016  
mov_i64 tmp2,$0x61b8  
ext16u_i64 tmp2,tmp2  
add_i64 tmp2,tmp2,cs_base  
ext32u_i64 tmp2,tmp2  
← mb $0x31  
qemu_ld_i64 tmp1,tmp2,leuw,2  
add_i64 tmp2,tmp2,$0x2  
ext32u_i64 tmp2,tmp2  
← mb $0x31  
qemu_ld_i64 tmp0,tmp2,leul,2  
and_i64 tmp0,tmp0,$0xffffffff  
st_i64 tmp0,env,$0x198  
st32_i64 tmp1,env,$0x1a0
```

```
set_label $L0
```

```
← exit_tb $0x280005f03
```



Does the TSOEnabler work as expected?

- Is TSO enforced on ARM?
- Sanity check with micro benchmark


How big is the performance impact of HW TSO?

- Comparison of software vs hardware fences
- PARSEC Benchmark Suite¹

¹ PARSEC Benchamrk Suite: <https://parsec.cs.princeton.edu/>


Evaluation: Does TSOEnabler work as expected?

TSO test program



Evaluation: Does TSOEnabler work as expected?

TSO test program



Count iterations over the whole array

Evaluation: Does TSOEnabler work as expected?

Experimental setup:

- Mac M1 8 Core CPU (3.2 GHz, 4 Performance, 4 Efficiency)
- AMD Ryzen 7 3800X 8 Core CPU (4.5 GHz, 16 Threads)

TSO test program until crash

	Max iterations	
ARM	≤1,600	repeated ~17k times
ARM + HW TSO	>10,000,000 - no reorder	Stopped after 15 mins
x86 AMD	>80,000,000 - no reorder	

TSO Enabler works ✓

Evaluation: Performance impact of hardware TSO

Experimental setup:

- Mac M1 8 Core CPU (3.2 GHz, 4 Performance, 4 Efficiency), 16 GB RAM
- PARSEC Benchmark Suite
- Arch Linux guest

List of builds

- SW Fences (Standard QEMU)
- HW Fences (usage of TSO)
- No Fences (incorrect)

2 different version for each build:


- “compile time”
QEMU build time
- “run time”
QEMU run time



Baseline: SW Fences (Standard QEMU) compile time

Evaluation: Performance impact of hardware TSO

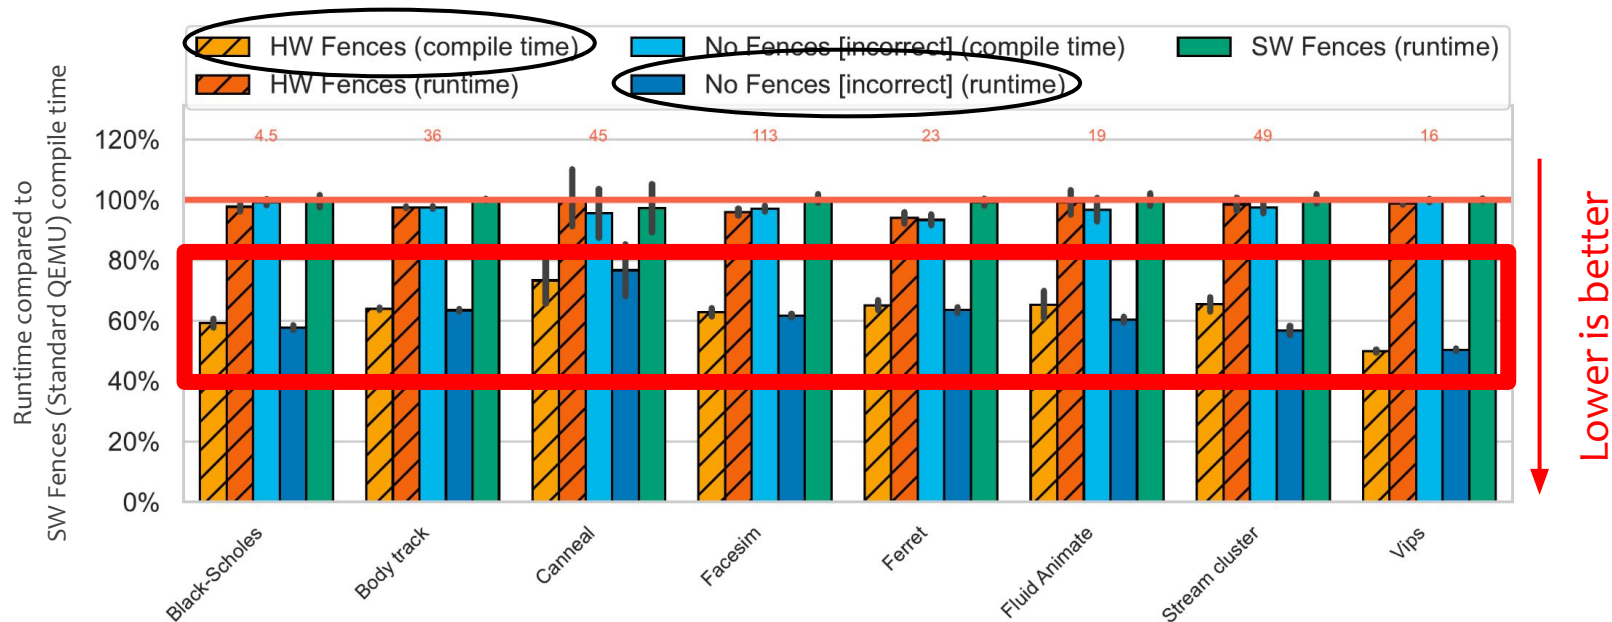
QEMU resources: 4 CPUs; 1 GB RAM



HW Fences \approx No Fences \Rightarrow HW TSO is cheap!

Evaluation: Performance impact of hardware TSO

QEMU resources: 4 CPUs; 1 GB RAM



Runtime variants \neq Compile time variants?


Investigation

- QEMU argument differences? **No**
- Benchmarking differences? **No**
- QEMU binaries swapped? **Sadly not**
- Thermal throttling? **Nope**
- Not using Performance Cores? **Also no**
- Background tasks? **No**
- CPU Feature (e.g. caching)? **No**

```
root@archlinux:/test/ ✓ > ./run_benchmarks.sh
Please enter a prefix for the results: No Fences
Benchmark iterations [20]:
Benchmark threads [4]:
Dataset [simlarge]:
```

```
Fence type: none, const: no
TSO support: no
```


```
Fence type: hardware, const: yes
TSO support: yes
Additional Info(m/v): d/d (x86: d)    Called tso_enable on thread 1272e9b1 with result 0
Called tso_enable on thread 118b29b1 with result 0
Called tso_enable on thread 118069b1 with result 0
Called tso_enable on thread 1198a9b1 with result 0
```



Ongoing Investigations

Impact between runs?

- 1 hour wait between runs
- one time 2 hour wait before a run



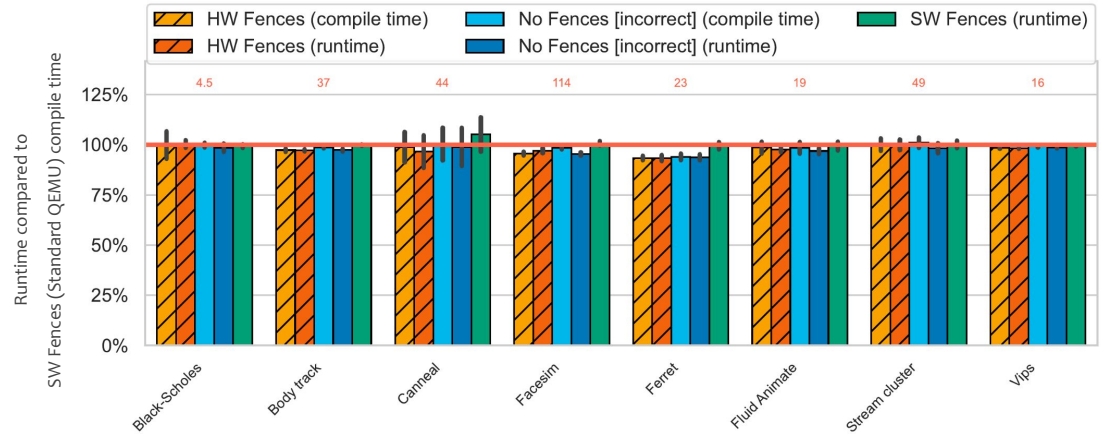
Ongoing Investigations

Impact between runs?

- 1 hour wait between runs
- one time 2 hour wait before a run
- run with cpu monitoring

Other options:

- Performance Profiler
- Different Mac M1



Summary

QEMU SW fences slow + incorrect

- Complex implementation
- performance on the slow end

QEMU + HW TSO:

- CPU responsible for memory order
- Simple implementation

BUT:

Performance benefit is unknown!

Try it out!

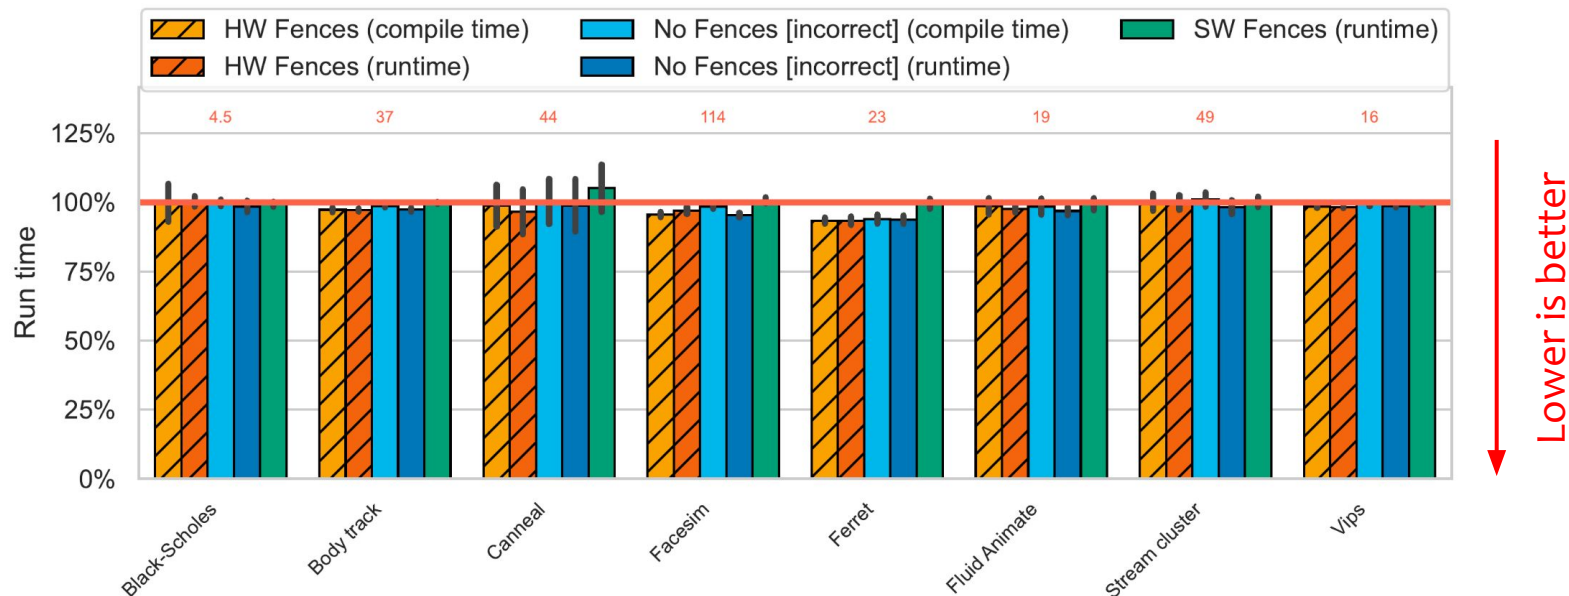
<https://github.com/fdevx/qemu>

Backup

Evaluation: Performance impact of hardware TSO

run times of the PARSEC benchmark suite

QEMU resources:
• 4 vCPU threads
• 4 GB RAM
PARSEC preset: “simlarge”



Everything is slow?

Generated instructions comparison

-- guest addr 0x00000000000efc9a	-- guest addr 0x00000000000efc9a		-- guest addr 0x00000000000efc9a
0x280006058: 91005295 add x21, x20, #0x14 (20)	0x280006058: 91005295 add x21, x20, #0x14 (20)	→	0x280006058: 91005295 add x21, x20, #0x14 (20)
0x28000605c: 2a1503f5 mov w21, w21	0x28000605c: 2a1503f5 mov w21, w21		0x28000605c: 2a1503f5 mov w21, w21
0x280006060: a97e0660 ldp x0, x1, [x19, #-32]	0x280006060: a97e0660 ldp x0, x1, [x19, #-32]		0x280006060: a97e0660 ldp x0, x1, [x19, #-32]
0x280006064: 8a551c00 and x0, x0, x21, lsr #7	0x280006064: 8a551c00 and x0, x0, x21, lsr #7		0x280006064: 8a551c00 and x0, x0, x21, lsr #7
0x280006068: 8b000021 add x1, x1, x0	0x280006068: 8b000021 add x1, x1, x0		0x280006068: 8b000021 add x1, x1, x0
0x28000606c: f9400020 ldr x0, [x1]	0x28000606c: f9400020 ldr x0, [x1]		0x28000606c: f9400020 ldr x0, [x1]
0x280006070: f9400c21 ldr x1, [x1, #24]	0x280006070: f9400c21 ldr x1, [x1, #24]		0x280006070: f9400c21 ldr x1, [x1, #24]
0x280006074: 91000ea3 add x3, x21, #0x3 (3)	0x280006074: 91000ea3 add x3, x21, #0x3 (3)		0x280006074: 91000ea3 add x3, x21, #0x3 (3)
0x280006078: 9274cc63 and x3, x3, #0xfffffffffffff000	0x280006078: 9274cc63 and x3, x3, #0xfffffffffffff000		0x280006078: 9274cc63 and x3, x3, #0xfffffffffffff000
0x28000607c: eb03001f cmp x0, x3	0x28000607c: eb03001f cmp x0, x3		0x28000607c: eb03001f cmp x0, x3
0x280006080: 54000b01 b.ne #+0x160 (addr 0x2800061e0)	0x280006080: 54000b01 b.ne #+0x160 (addr 0x2800061e0)		0x280006080: 54000b01 b.ne #+0x160 (addr 0x2800061e0)
0x280006084: b8756835 ldr w21, [x1, x21]	0x280006084: b8756835 ldr w21, [x1, x21]		0x280006084: b8756835 ldr w21, [x1, x21]
0x280006088: f9000275 str x21, [x19]	0x280006088: f9000275 str x21, [x19]		0x280006088: f9000275 str x21, [x19]
			0x28000608c: d50339bf dmb ishld
			0x2800060f4: a97e0660 ldp x0, x1, [x19, #-32]
			0x2800060f8: 8a551c00 and x0, x0, x21, lsr #7
			0x2800060fc: 8b000021 add x1, x1, x0
			0x280006100: f9400020 ldr x0, [x1]
			0x280006104: f9400c21 ldr x1, [x1, #24]
			0x280006108: 91000ea3 add x3, x21, #0x3 (3)
			0x28000610c: 9274cc63 and x3, x3, #0xfffffffffffff000
			0x280006110: eb03001f cmp x0, x3
			0x280006114: 54000b01 b.ne #+0x160 (addr 0x280006280)
			0x280006118: b8756835 ldr w21, [x1, x21]
			0x28000611c: f9000275 str x21, [x19]