

Bachelor's thesis

Programming model for hybrid persistent memory systems

Matthias Werndle


Advisor: Dimitrios Stavrakakis

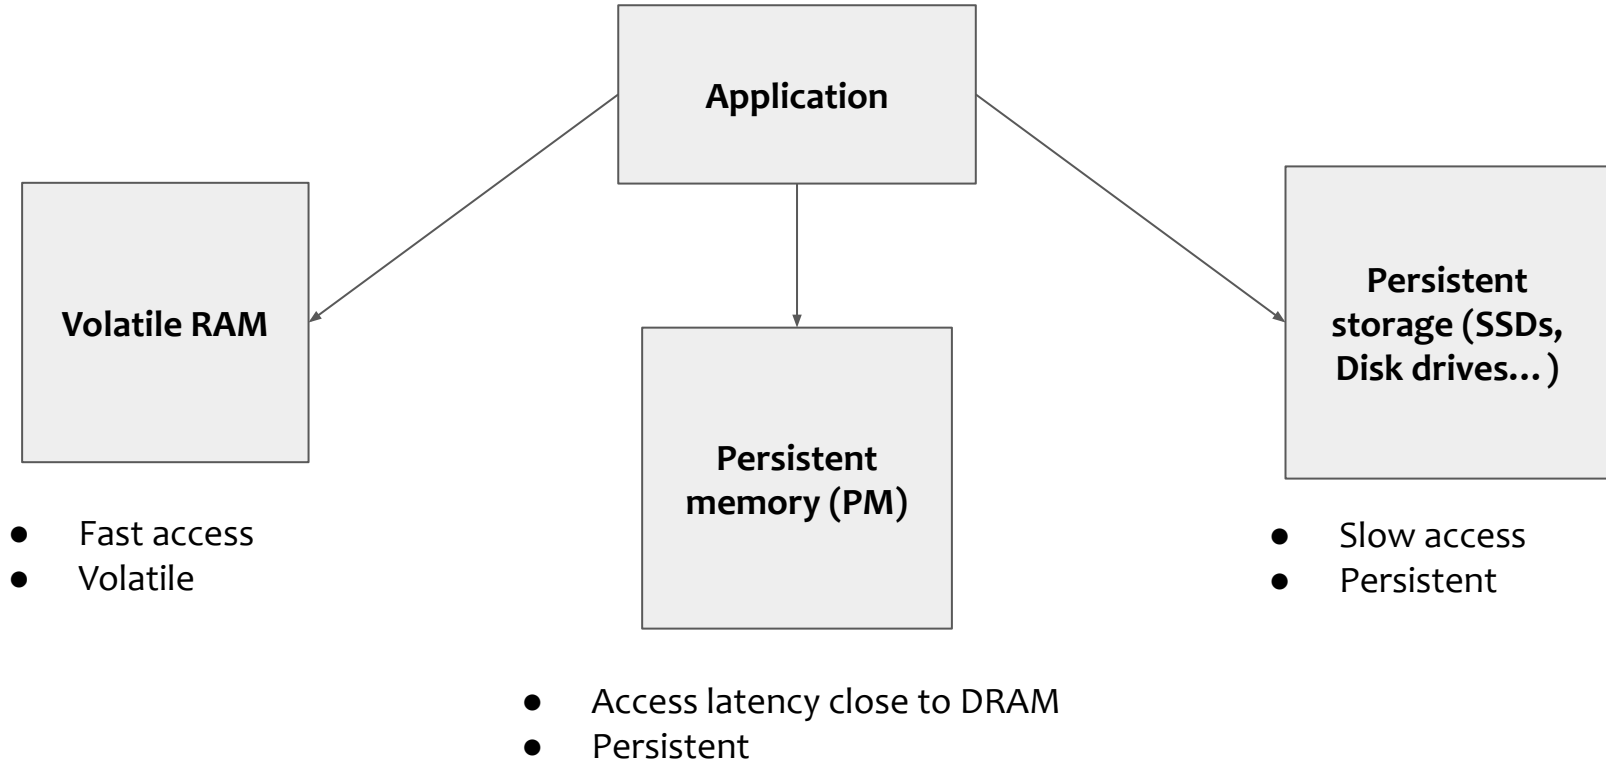
Chair of Decentralized Systems Engineering

<https://dse.in.tum.de/>



15.04.2022 – 16.08.2022

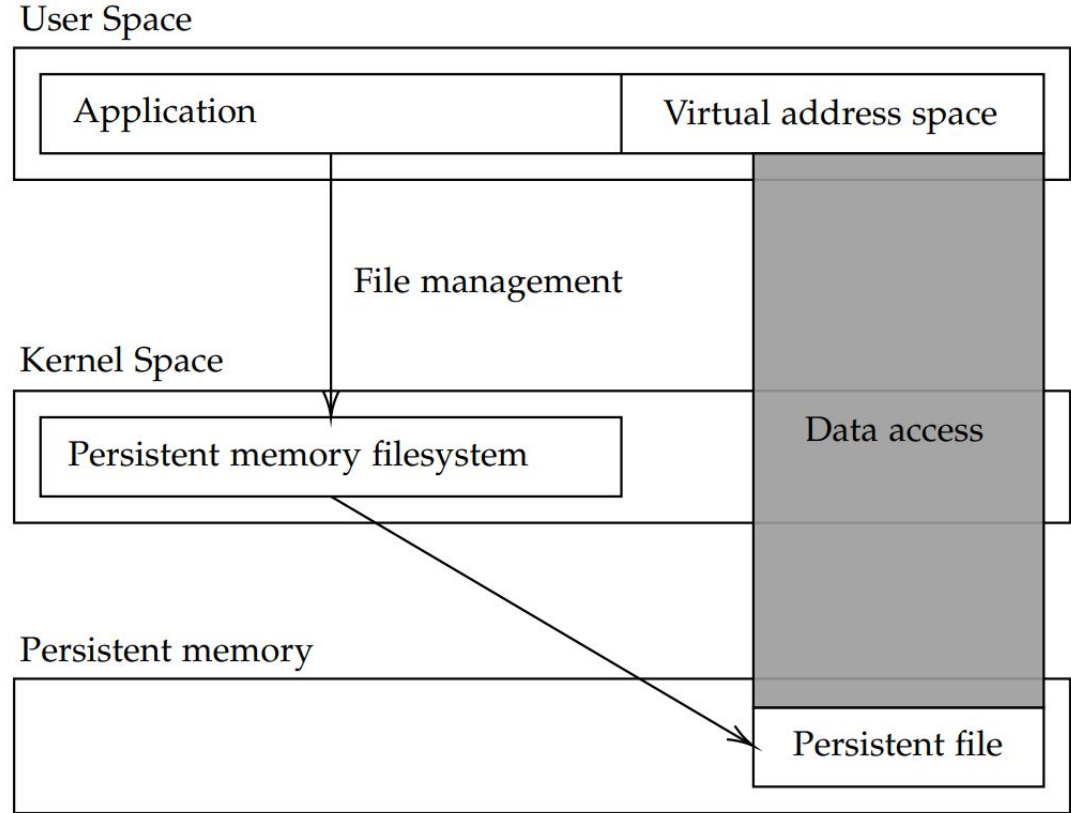




Background - Programming model

- PM space is structured as files or regions
- Kernel handles PM file access
- Memory mapping allows direct access to PM file


➡ **Pointer relocation requires special handling**



Two pointer types:

- Volatile Memory (**VM**) Pointers
- Persistent Memory (**PM**) Pointers


PM Pointer layout:



- PM pointer is **independent** of memory mapped address space
- Requires conversion from **PM** pointer to **VM** pointer to access data
- Requires **additional** data types


- ➡ **Different data structures** are required for PM support depending on the used library
- ➡ It is hard to **convert** existing **VM** applications into **PM** applications
- ➡ Programs are **depending** heavily on one library

Design - Overview

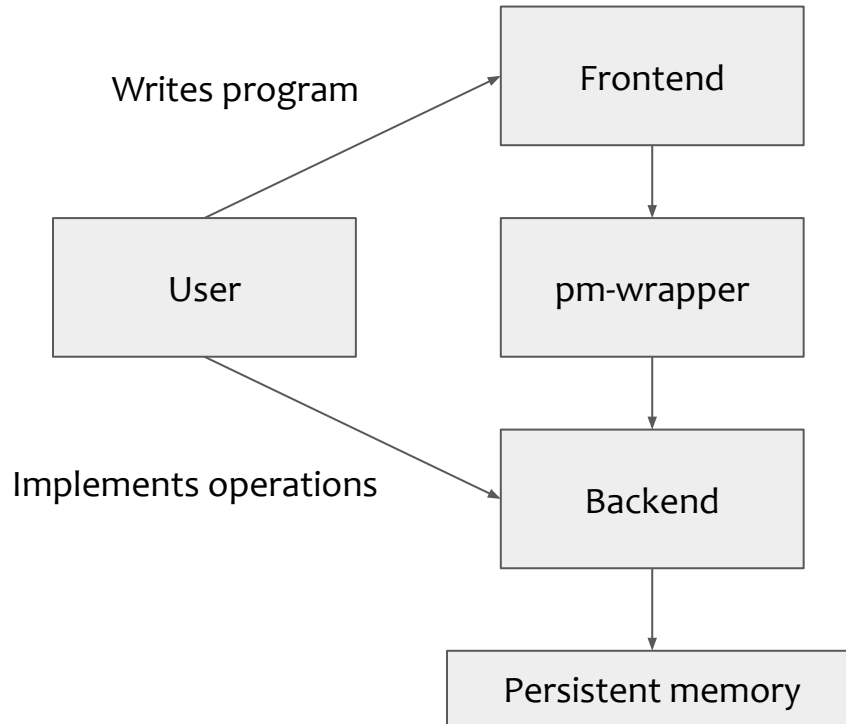


→ Allows pointer access to persistent memory like with C pointers

Design - Overview



- Allows transparent c like pointer access to persistent memory
- Layer which consumes source code to perform PM operations



- Allows transparent c like pointer access to persistent memory
- Layer which consumes source code to perform PM operations
- Manages connection to PM
- Implementation can use library such as PMDK

- Introducing a new implicit type system to differentiate between pointer types

```
1 // VM pointer
2 int* vm_ptr = malloc(...);
3
4 // PM pointer
5 int* pm_ptr = pm_alloc(...);
```


- Usual pointer operations can be performed on PM pointers
- Undetectable pointer types have to use attributes



Overall frontend allows PM access similar to VM programs


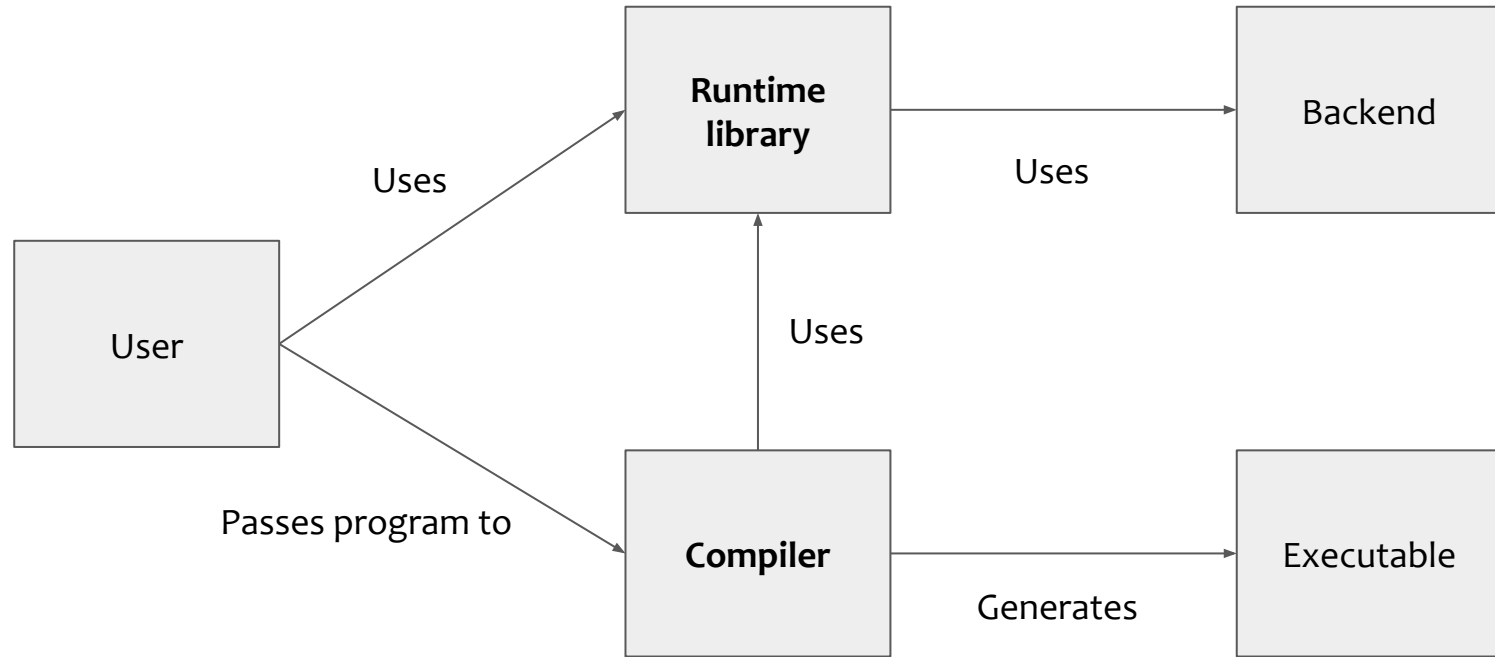
Design - Backend

- Handling low level PM access
- Operations types:



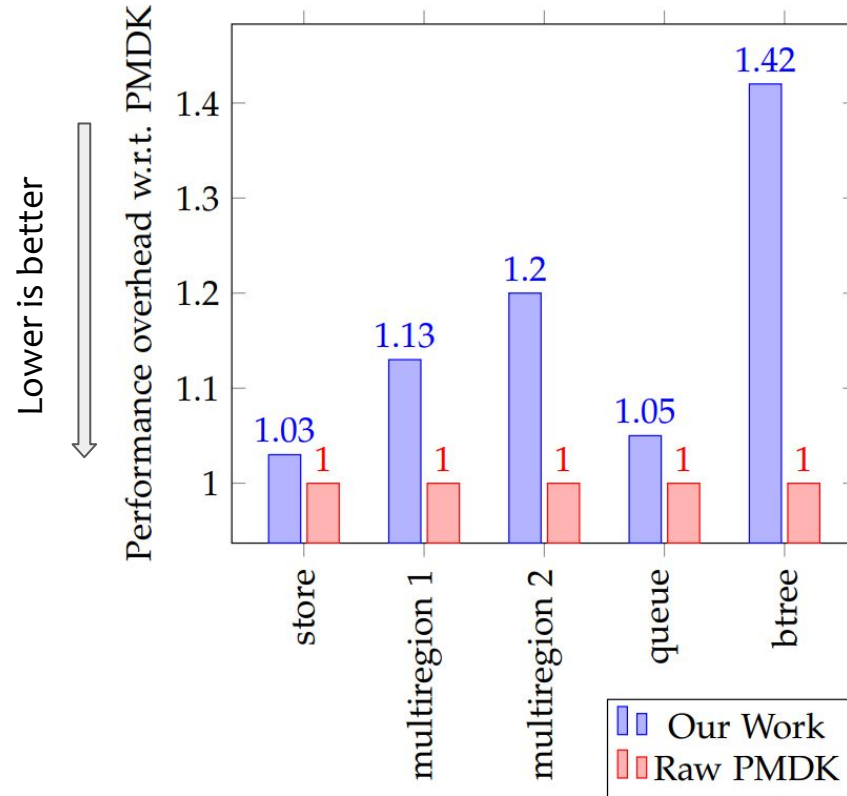
Backend allows to switch implementations without changing the frontend

Implementation - Overview



Two major parts: **Runtime library** and **Compiler**

Evaluation - Runtime overhead



- Compares raw PMDK versions with pm-wrapper versions implemented with PMDK
- Compiled with Clang-13 and O3 optimization



Overall low performance overhead



In some cases compiler inserts have to be optimized (btree)

Migration of a VM hashtable to a PM hashtable

Change	Description	Count
Function relacements	Switching a called function with another	6
Structure changes	Adding, removing or altering members of structs	0
Additional logic	Inserting additional required logic	4
PMI type attributes	Required type attributes	5
Pointer operation	Required pointer operations changes	0



Migration is improved, but in some cases the compiler still requires type attributes

- Quicker migration to PM support
- Simpler migration between PM libraries
- Performance overhead is low



pm-wrapper provides an additional layer to improve PM programming



Source Code: <https://github.com/Teppichseite/pm-wrapper>

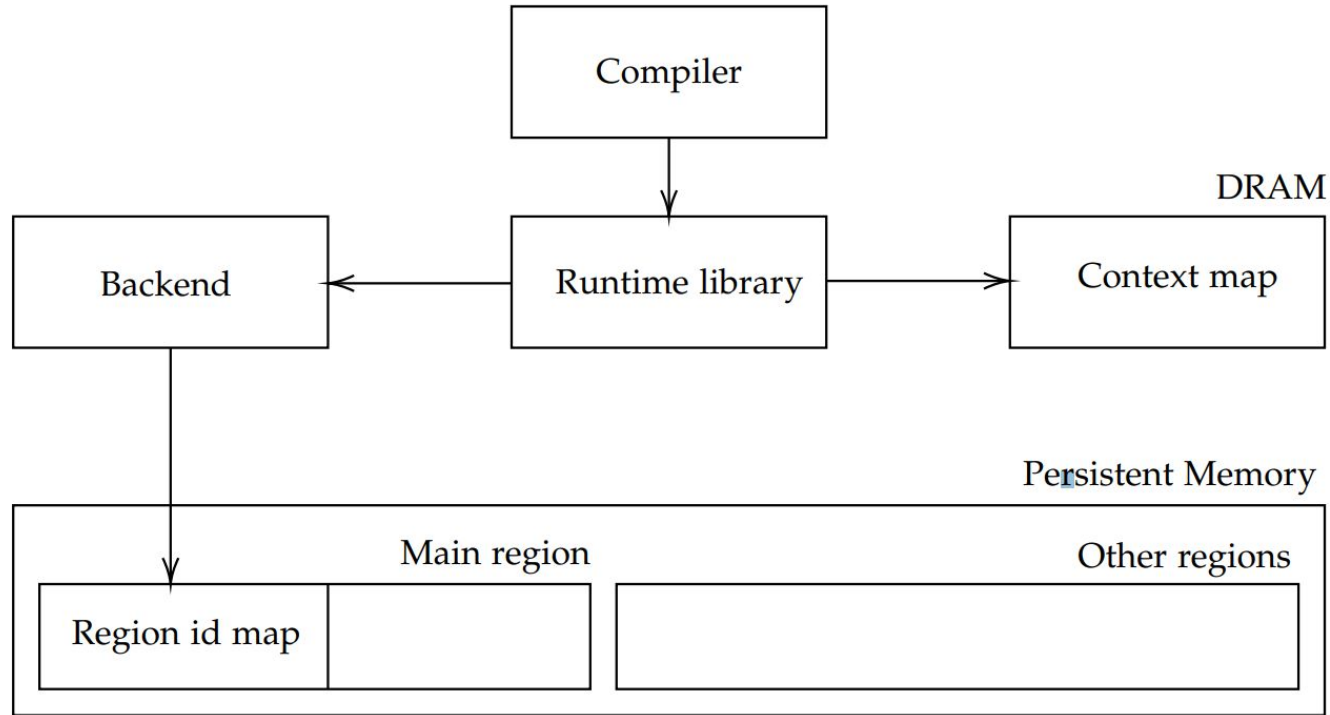
Backup

Design - Frontend example

```
1 struct Store {
2     int *ptr_a;
3     int *ptr_b;
4 };
5
6 struct Store *store1 = (struct Store *)pm_alloc(sizeof(struct Store));
7
8 // PMI type PM
9 int *ptr1 = store1->ptr_a;
10
11 struct Store *store2 = (struct Store *)malloc(sizeof(struct Store));
12
13 // PMI type PM by attribute
14 PM int *ptr2 = store2->ptr_a;
15
16 // PMI type VM
17 int *ptr3 = store2->ptr_b;
```


```
1 struct PmBackend {
2     int (*init)();
3     int (*open_or_create)(PmBackendContext *context, bool *created_new);
4     void (*close)(PmBackendContext *context);
5     void (*finalize)();
6     pm_region_offset (*get_root)(PmBackendContext *context);
7     pm_region_offset (*alloc)(PmBackendContext *context, size_t size);
8     pm_region_offset (*calloc)(PmBackendContext *context, size_t size);
9     void (*free)(PmBackendContext *context, pm_region_offset offset);
10    void (*read_object)(PmBackendContext *context, pm_region_offset offset);
11    void (*write_object)(PmBackendContext *context, ...);
12};
```

Implementation - Runtime



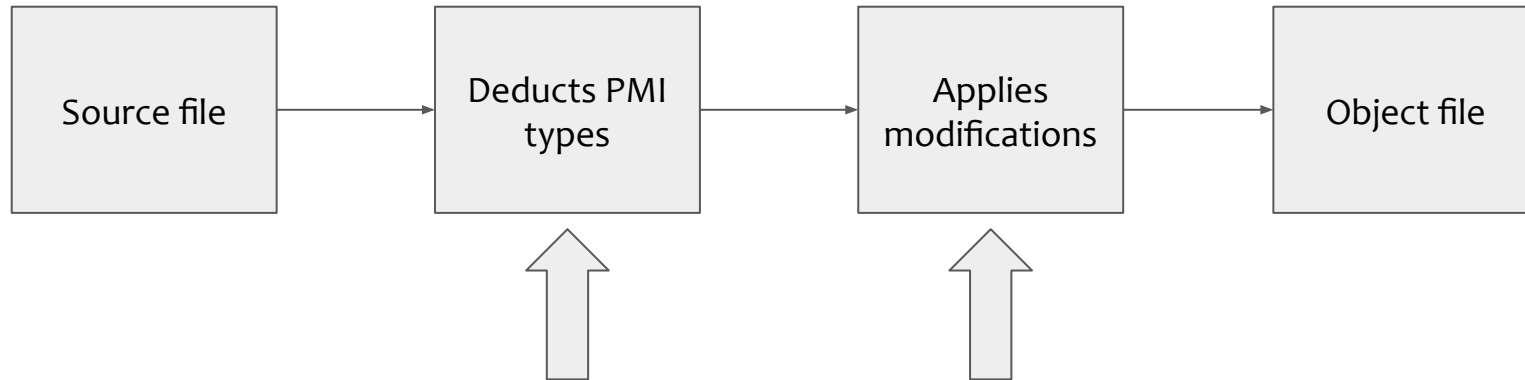
Implementation - Runtime

- Offers **region** and **memory management** interface
- Handles **mapping** between **frontend pointers** and backend **persistent pointers**
- Mapping can be realized in two ways:
 - Hash map
 - Fixed size array lookup



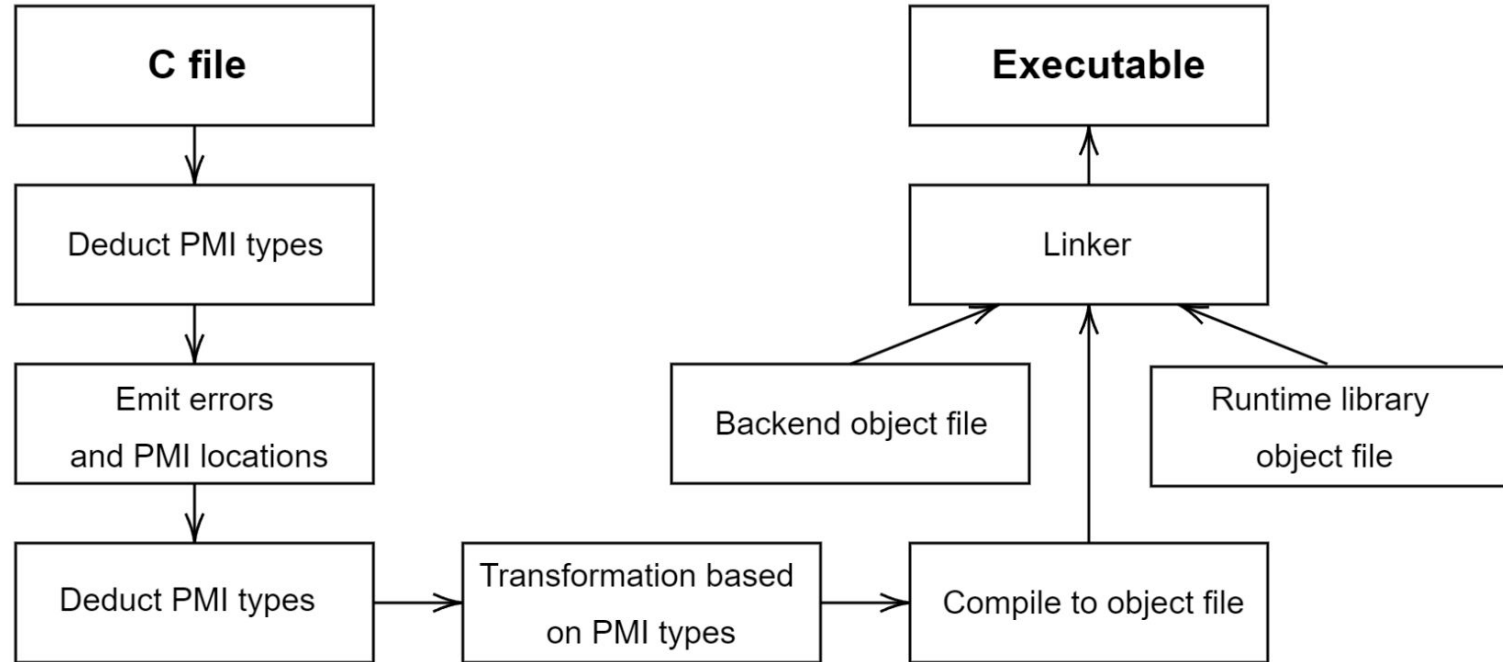
Implementation - Compiler

- Implemented with Clang LibTooling
- Analyzes directly the C source code
- Deducts pointer types
- Rewrites a copy of the actual source code



Main implementation effort for pm-wrapper

Implementation - Compiler chain



Backend implementation

```
1 void write_object(void *dst, char *data, size_t len) {  
2     // Perform logging  
3     pmemobj_tx_add_range(dst, 0, len);  
4     // Copy data  
5     pmemobj_memcpy(dst, data, len);  
6 }
```

Frontend usage

```
1 PM struct Data *data;  
2 pm_tx_start();  
3 data->a = 1;  
4 data->b = 2;  
5 pm_tx_end();
```