

Coyote

A Comprehensive Study of Coyote Memory System

Yuan Luo

Supervisor: Atsushi Koshiba, Jiyang Chen

Chair of Distributed Systems and Operating Systems

<https://dse.in.tum.de/>



- Sharing FPGA between different user logics is difficult due to potential I/O conflicts
- Programming FPGA applications targeting different I/O devices remains challenging due to insufficient support
- An easy-to-use FPGA I/O acceleration framework is needed (**FAIR**)

- AmorphOS
 - Focuses only on memory I/O
- Coyote
 - Supports memory I/O and network I/O
 - Lacks efficient I/O transaction scheduling

We need a smarter FPGA I/O scheduler for different I/O devices

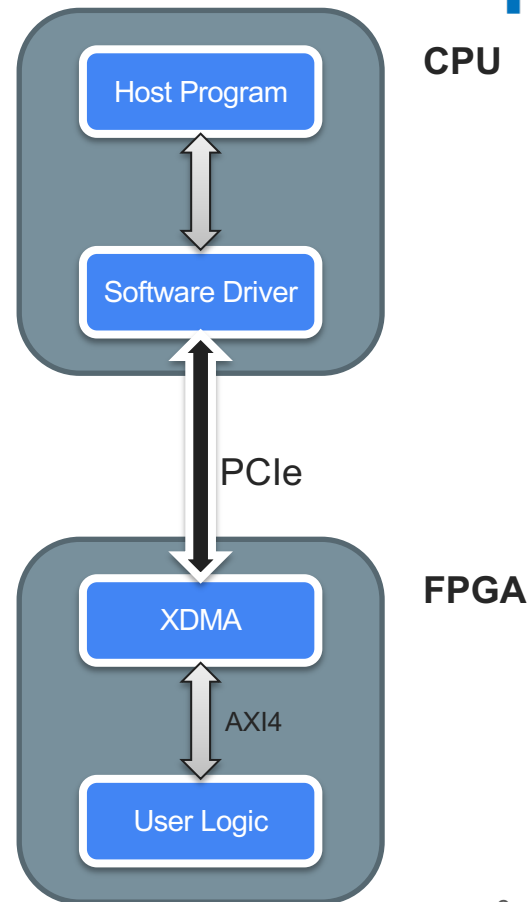
A survey on Coyote to implement FAIR I/O scheduler is required

- How Coyote implements virtual FPGA and virtual memory at the hardware level
- How Coyote memory system works
- How to integrate an I/O scheduler into Coyote memory system

- Coyote Architecture
- Coyote Memory System
- Debugging Coyote
- Evaluation

Communication between CPU and FPGA

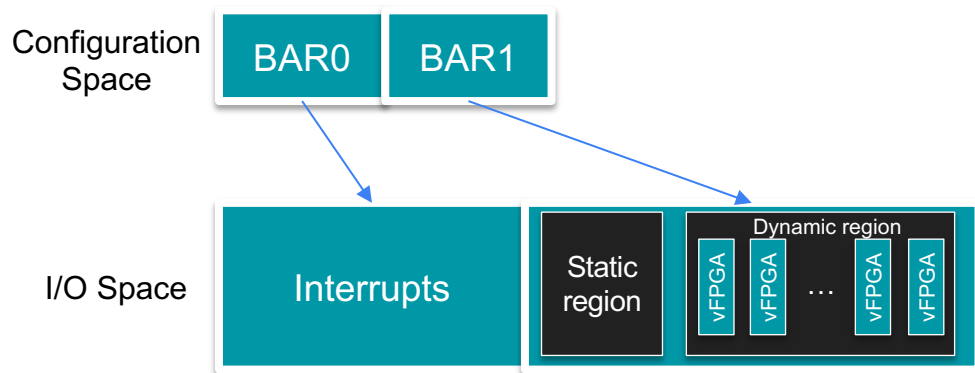
- Driver connects host program with PCIe
- XDMA configures PCIe and communicates with user logic



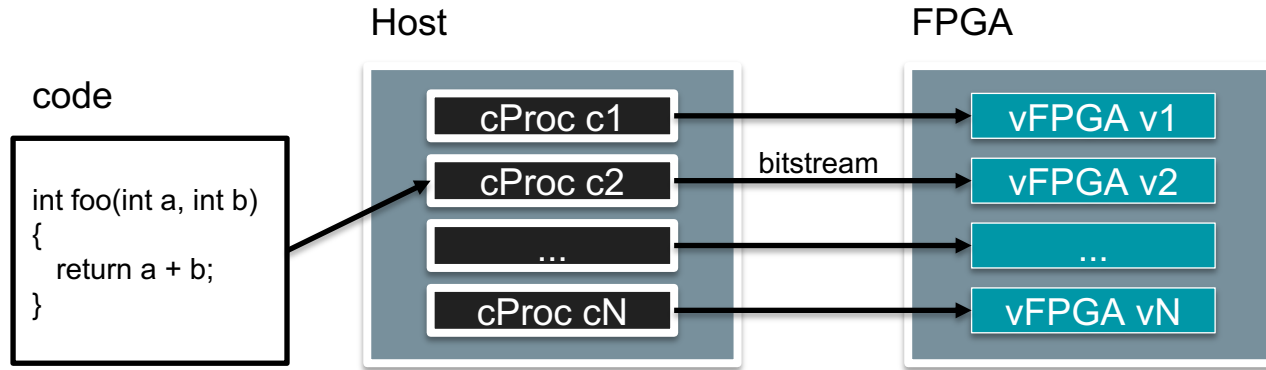
Static region and dynamic region

PCIe memory space contains three address spaces:

- I/O address space
- Memory address space
- Configuration address space



Coyote Process (cProc) and vFPGA



- Program code is first synthesized into bitstream
- cProc loads the bitstream and reconfigures vFPGA

- ~~— Coyote Architecture~~
- Coyote Memory System
- Debugging Coyote
- Evaluation

Virtual address

- A strategy similar to CUDA
- The page size can be either 4 KB or 2 MB


Code

```
void* hMem = cproc.getMem({CoyoteAlloc::HUGE_2M, n_pages});
```

Dynamic
region of a
vFPGA in
driver



vFPGA
wrapper
on FPGA

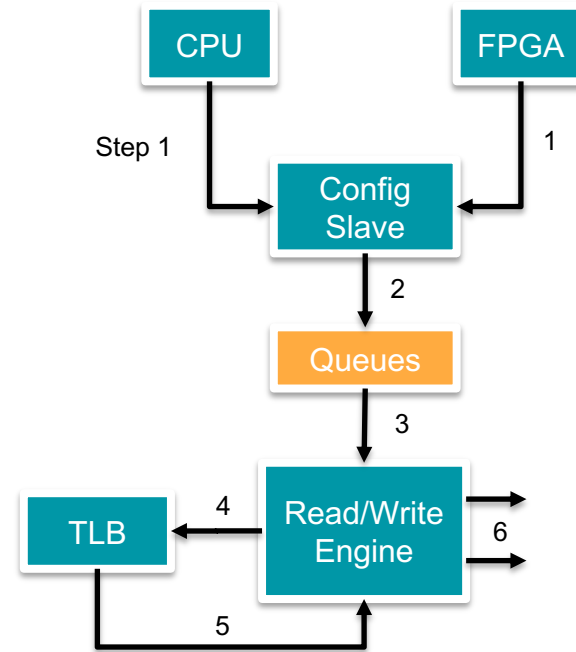


Memory Manage Unit (MMU):

- Config Slave
- Queue
- Read/Write engines
- ...

Memory management

1. Collect memory access requests
2. Process and distribute requests
3. Forward memory access requests
4. Refer to TLB for possible address translation
5. Return physical address or issue error
6. Issue memory access



Outline




- ~~— Coyote Architecture~~
- ~~— Coyote Memory System~~
- ~~— Debugging Coyote (skipped)~~
- Evaluation

- Host-initiated data transfer

- A loopback data transfer between host and FPGA

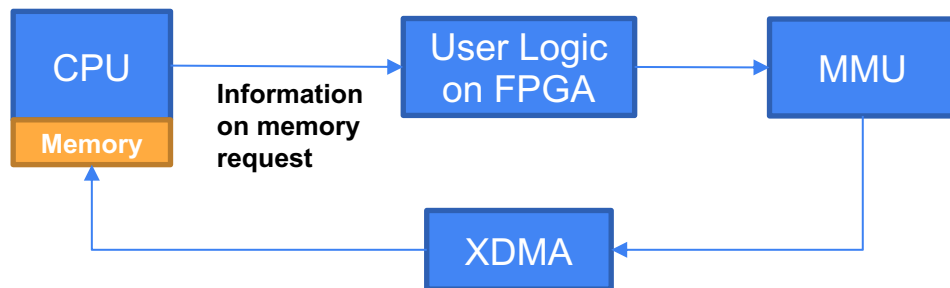
Test Platform (Sakura)

- CPU: Intel(R) Xeon(R) Gold 6238R
- RAM: 256 GB
- FPGA: Alveo U50 (8GB ROM)



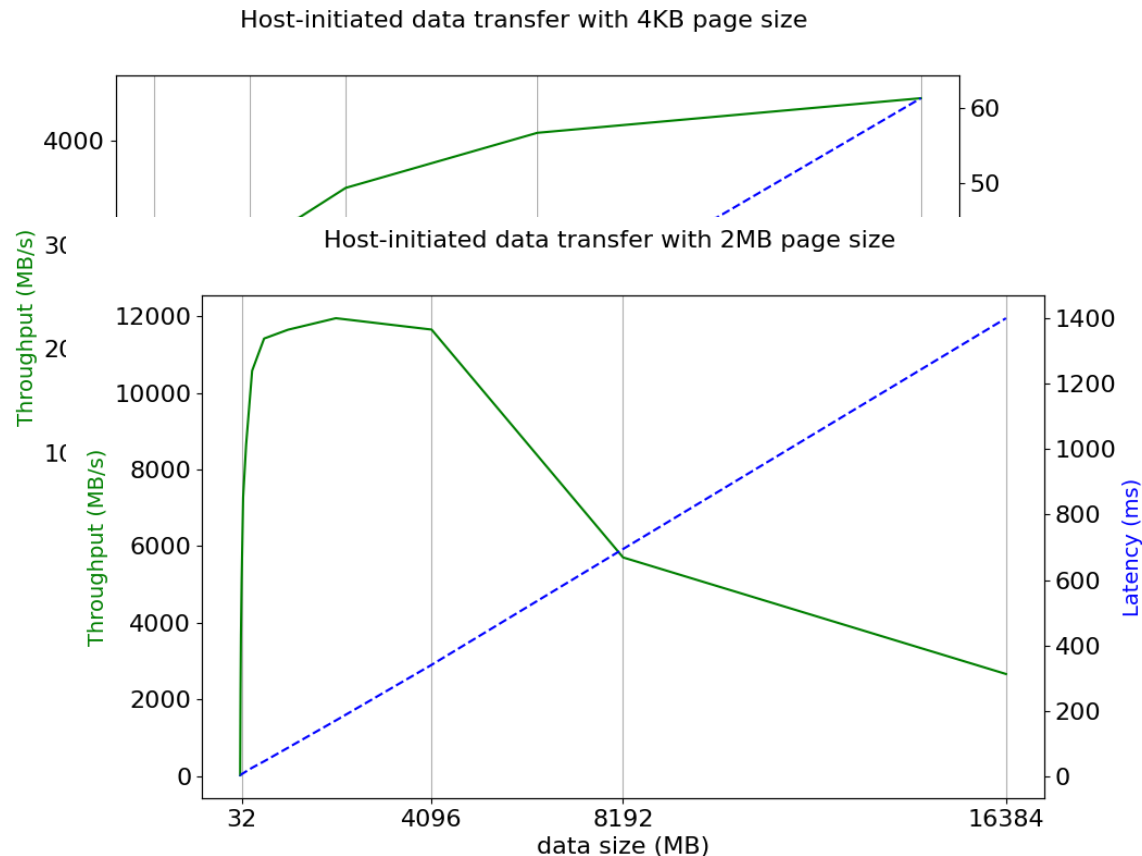
- FPGA-initiated data transfer

- FPGA tries to access memory on the host




Host-initiated data transfer

- Higher Throughput as payload's size grows
- Alveo U50 has only 8 GB Memory



- Write operation performs better than read operation
- Higher priority for I/O scheduling



- Coyote uses PCIe address spaces to enable host-FPGA communication and process abstraction.
- The memory system enables virtual addresses for FPGA memory on the host side and implements the memory management unit (MMU) for each vFPGA
- A possible priority-based I/O scheduler can be implemented on top of memory access descriptor queues