

Arancini

Emulating x86 binaries on the RISC-V architecture

Simon Kammermeier

Advisor: Dr. Redha Gouciem

Chair of Distributed Systems and Operating Systems

<https://dse.in.tum.de/>



15.11.2022 – 15.04.2023

Reason for binary translation

- Program compatibility key requirement for adoption of new CPU architectures
- Recompilation often not possible
- General solution: Translate binary to new CPU architecture

Existing solutions



- QEMU: Open source, support for many ISAs, translates at execution time
- Lasagne, mctoll: Open source, x86-64 via LLVM to ARM, translates ahead of time
- Apple Rosetta2: Proprietary, x86-64 to ARM, exploits special hardware, combines translation ahead of time and at execution time

Approaches of Binary Translation

Static binary translation

Dynamic binary translation

Ahead-of-time

At Runtime

Slow, complex translation possible

Fast, simple translations

Optimized code

Unoptimized code

Problem: Code Discovery

⇒ Combine both approaches to get fast runtime and full coverage

Hybrid binary translation

System design goals:

- Modularity
- General hardware support
- Multithreaded application support

Implementation of the dynamic path for x86 to RISC-V

Outline

- ~~Motivation~~
- Background
 - RISC-V
- Design
- Implementation
- Evaluation

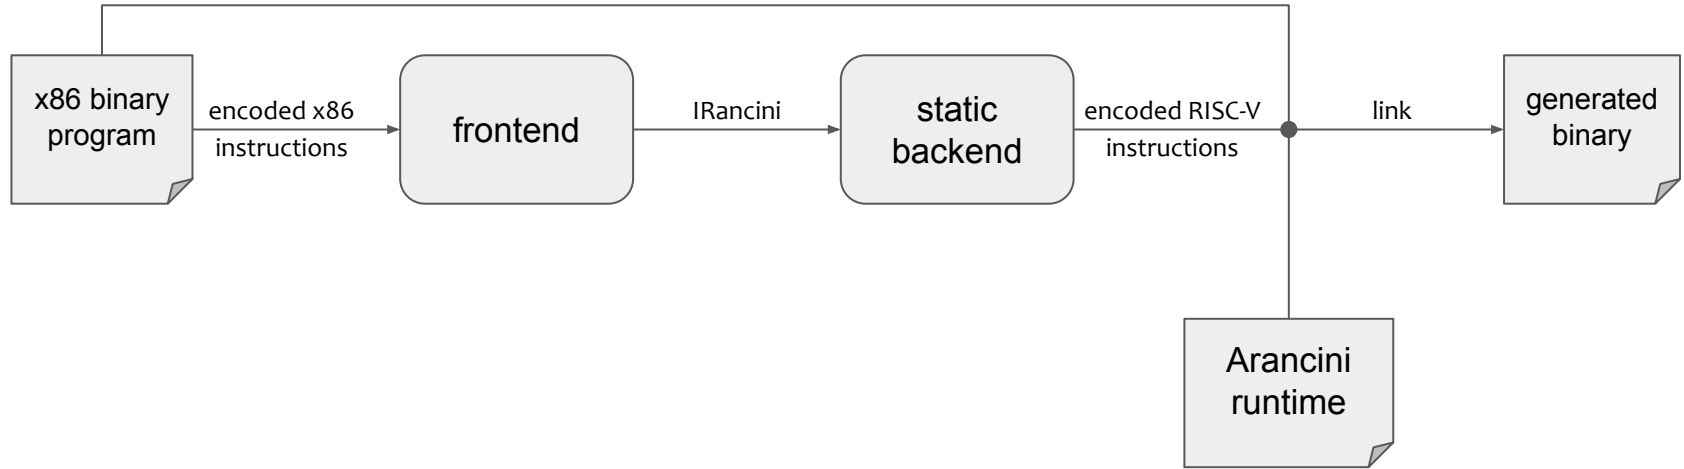
Background: RISC-V

- Initially research project at UC Berkeley
- Goal: free, open source ISA for real world implementations
- Key design differences to x86:
 - RISC philosophy
 - Load-Store Architecture
 - No condition flags

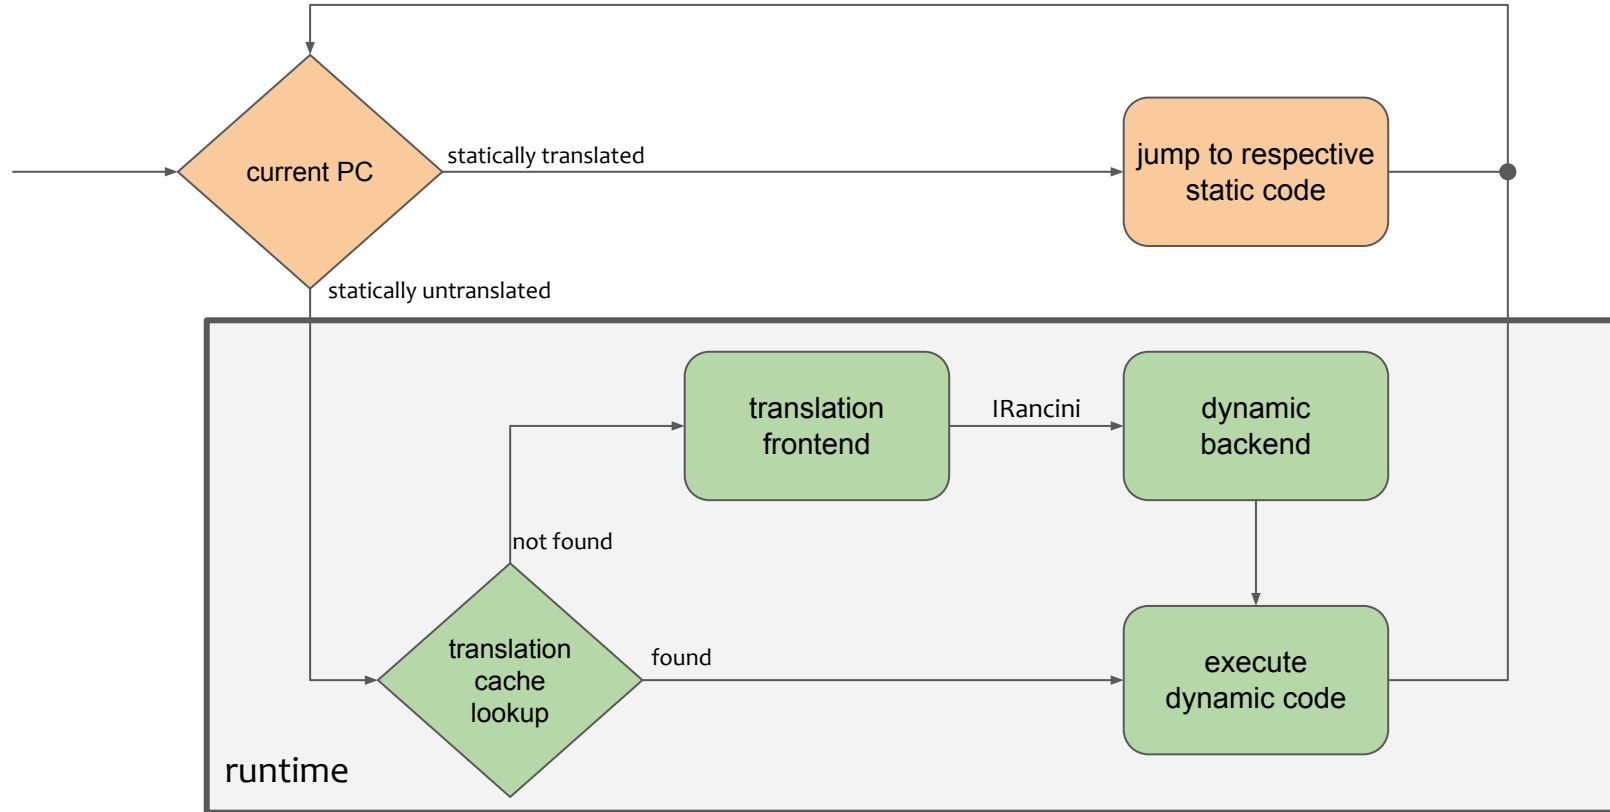
Outline

- ~~Motivation~~
- ~~Background~~
- Design
 - System overview
 - IRancini
- Implementation
- Evaluation

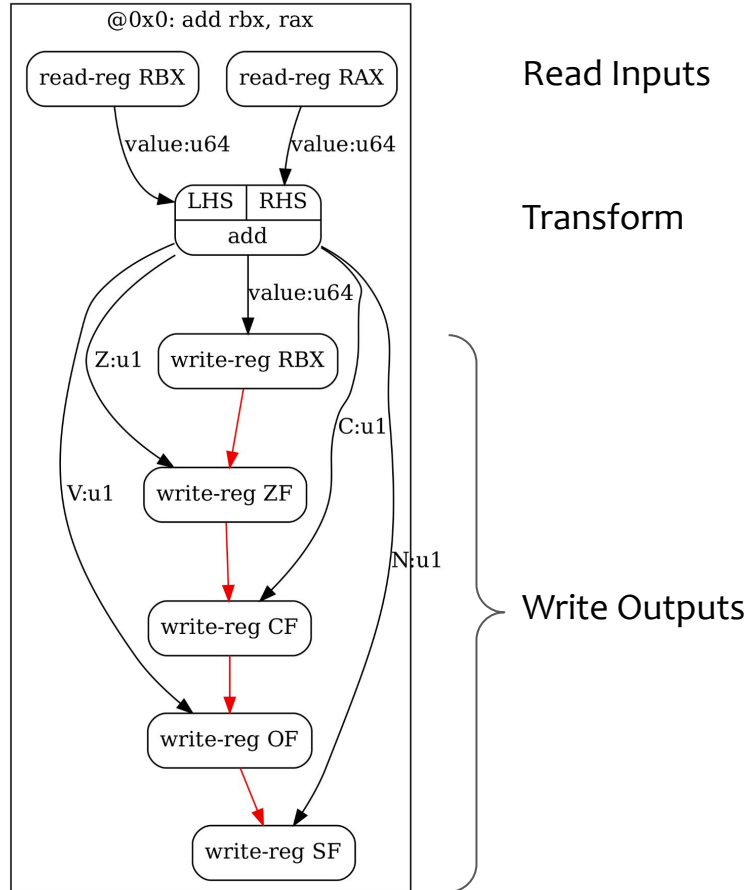
System overview: Static Phase



System overview: Runtime Phase



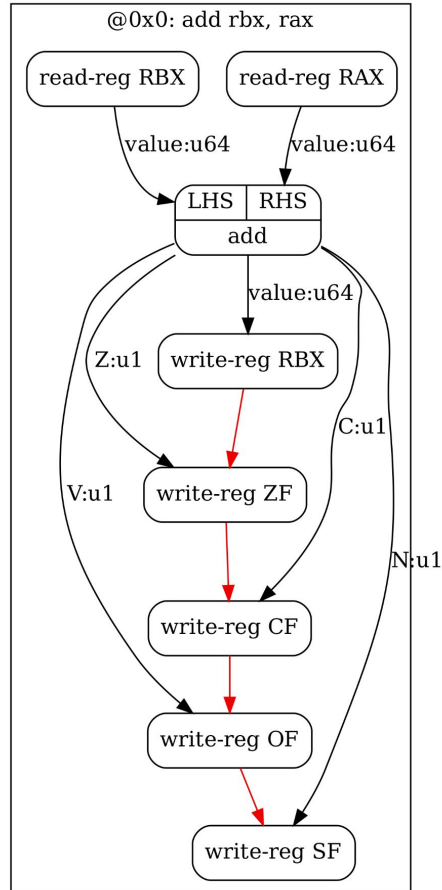
Design overview: IRancini



Outline

- ~~Motivation~~
- ~~Background~~
- ~~Design~~
- Implementation
- Evaluation

Implementation



```
ld    a0, 32(fp)
```

```
ld    a1, 8(fp)
```

```
add   s1, a0, a1
```

```
sltz  s9, a0
```

```
slt   s10, s1, a1
```

```
xor   s10, s10, s9
```

```
sltu  s9, s1, a1
```

```
seqz  s8, s1
```

```
sltz  s11, s1
```

} Overflow flag

```
sd    s1, 32(fp)
```

```
sb    s8, 392(fp)
```

```
sb    s9, 393(fp)
```

```
sb    s10, 394(fp)
```

```
sb    s11, 395(fp)
```

Outline

- ~~Motivation~~
- ~~Background~~
- ~~Design~~
- ~~Implementation~~
- Evaluation


- Are the programs emulated correctly?
- How much overhead is created by the translation?

- Experimental setup:
 - HiFive Unmatched RISC-V Development Kit by SiFive
 - U74-MC CPU:
 - 4 U74 cores: RV64GC
- Test programs:
 - Various single threaded integer benchmarks provided as part of phoenix
 - Comparison to benchmarks natively compiled

Correctness

- No crashes in any of the test programs
- Comparison of benchmark outputs showed no differences
- No strict proof of correctness but good evidence

Runtime overhead



Optimization opportunities

- Register mapping to avoid memory accesses
- Elimination of redundant instructions
 - unused flag computations
- Block chaining

Summary

Translation approach of split front and backend very feasible for x86 to RISC-V

Correct translation of single threaded programs using integer scalar instructions

Performance is currently not good

Some easy optimizations could improve performance

Backup

Design overview: IRancini

